



# PNEUMONIA DETECTOR



+

1

# WHAT IS PNEUMONIA?



+

# ABOUT THE DISEASE

Pneumonia is an infection that affects the air sacs in one or both lungs. The air sacs fill with fluids causing difficulty with breathing as well as coughing, fever, and chills.

The most common mode of evaluation and diagnosis for the infection is by an x-ray exam. X-ray images use varying colors to display density; patients with pneumonia will have white spots (called infiltrates) in their x-ray scans.



# OUR MOTIVATION

1.5 million people were diagnosed with pneumonia back in 2018, and nearly 38.8% of those cases were misdiagnosed. Pneumonia is the 2nd most misdiagnosed condition in the US.

## MISDIAGNOSIS

Misdiagnosis often occurs due to incorrect readings of x-ray images. The blacks, grays, and whites, make it extremely hard to spot the white clouds that aid in a physician's diagnosis from an x-ray image.

When undiagnosed/ untreated, the infection can lead to respiratory failure and even death.

## MOTIVATION

Our project wanted to test AI accuracy in analyzing x-ray imaging to spot and diagnose pneumonia. A pneumonia detecting AI system, would help perform tasks in reading x-ray imaging. If successful, the project/ AI could decrease the number of misdiagnosis and help save hundreds of lives!



+

2



+

# Understanding and Visualizing Our Data

Making the foundation of our  
machine

# Data Columns

01

## Class

Tells us whether the image shows healthy lungs or lungs with pneumonia (represented as either 0 or 1, binary classification)

02

## Split

Tells us whether the data is from our training set or testing set

03

## Index

Tells us where in the data set the image is located

|      | class | split | index |
|------|-------|-------|-------|
| 0    | 0.0   | train | 0     |
| 1    | 0.0   | train | 1     |
| 2    | 1.0   | train | 2     |
| 3    | 0.0   | train | 3     |
| 4    | 1.0   | train | 4     |
| ...  | ...   | ...   | ...   |
| 2395 | 1.0   | test  | 2395  |
| 2396 | 0.0   | test  | 2396  |
| 2397 | 0.0   | test  | 2397  |
| 2398 | 1.0   | test  | 2398  |
| 2399 | 0.0   | test  | 2399  |

2400 rows x 3 columns

# Training Vs Testing Data

## Training Data

- Training Manual
- Helps Machine learn its task
- Much larger than testing data
- Useful in teaching data how to detect patterns
- "Homework" for the machine


## Testing Data

- Test for the machine
- Sees if the machine can recognize patterns outside of the training set
- Problems machine has not seen before
- Is the machine ready to be put to use?

Make Sure Your Data Is Not Overfit!!!!

# What Does Our Data Look Like?

## Plotting Our Data






# Manipulating Our Data

```
for i in range(20):  
    for j in range(10):  
        rect_image[i,j,0] = 1  
plot_one_image(rect_image)
```

## Highlighting certain areas in image

- Rect\_image[i,j,0] = 1
  - I = x value of image (what should the width and x-position of rectangle be?)
  - J = y value of image (what should the length and y-position of the rectangle be?)
  - By setting pixel = 1, making the pixel fully red



## Important uses of function

- Not all of image is necessary for detection
- Zoning in on specific locations makes it easier to detect
- Being able to compare those locations to others
- Makes it easier to detect specific patterns in data



+


3

# Creating and Applying Neural Networks



+

# Neural Networks



Neural Networks look something like this



The neural network is built using the packages known as 'tensorflow' and 'keras'

## Sample created and tested neural network

```
#YOUR CODE HERE
from sklearn.neural_network import MLPClassifier

(train_data, train_labels) = get_train_data(flatten = True)
(test_data, test_labels) = get_test_data(flatten = True)
mlp = MLPClassifier(hidden_layer_sizes = 5)
mlp.fit(train_data, train_labels)
storage = mlp.predict(test_data)
print(accuracy_score(test_labels, storage))
#YOUR CODE HERE
```


0.5

# Applying Neural Networks to Medical Imaging

- Our model was calibrated through the 'relu' and 'sigmoid' function
- We are given 'images' of shape '(64,64,3)', each assigned a label PNEUMONIA or HEALTHY.

```
dense = DenseClassifier(hidden_layer_sizes = (64,32))
```

```
cnn = CNNClassifier(num_hidden_layers = 1)
```



# Fitting

```
model_history = model.fit(train_data, train_labels, epochs = 100, validation_data = (test_data, test_labels), shuffle = True, callbacks = [monitor])
```

## Our Model

```
model_history1 = dense.fit(train_data, train_labels, epochs = 100, validation_data = (test_data, test_labels), shuffle = True, callbacks = [monitor])
model_history2 = cnn.fit(train_data, train_labels, epochs = 100, validation_data = (test_data, test_labels), shuffle = True, callbacks = [monitor])
```

```
Epoch 1/100
63/63 [=====] - 0s 5ms/step - loss: 0.2739 - accuracy: 0.8950 - val_loss: 0.5445 - val_accuracy: 0.7250
Epoch 2/100
63/63 [=====] - 0s 4ms/step - loss: 0.2523 - accuracy: 0.9015 - val_loss: 0.6019 - val_accuracy: 0.7220
Epoch 3/100
63/63 [=====] - 0s 5ms/step - loss: 0.2518 - accuracy: 0.9045 - val_loss: 0.5639 - val_accuracy: 0.7150
Epoch 4/100
63/63 [=====] - 0s 5ms/step - loss: 0.2574 - accuracy: 0.9040 - val_loss: 0.5502 - val_accuracy: 0.7220
Epoch 5/100
63/63 [=====] - 0s 5ms/step - loss: 0.2604 - accuracy: 0.9040 - val_loss: 0.6377 - val_accuracy: 0.7200
Epoch 6/100
63/63 [=====] - 0s 5ms/step - loss: 0.2578 - accuracy: 0.8970 - val_loss: 0.5987 - val_accuracy: 0.7250
Epoch 7/100
63/63 [=====] - 0s 5ms/step - loss: 0.2563 - accuracy: 0.9020 - val_loss: 0.5533 - val_accuracy: 0.7100
Epoch 8/100
63/63 [=====] - 0s 5ms/step - loss: 0.2587 - accuracy: 0.8960 - val_loss: 0.6787 - val_accuracy: 0.7120
Epoch 9/100
63/63 [=====] - 0s 5ms/step - loss: 0.2444 - accuracy: 0.9000 - val_loss: 0.6323 - val_accuracy: 0.7170
Epoch 10/100
63/63 [=====] - 0s 5ms/step - loss: 0.2515 - accuracy: 0.9040 - val_loss: 0.6306 - val_accuracy: 0.7050
Epoch 11/100
63/63 [=====] - 0s 5ms/step - loss: 0.2259 - accuracy: 0.9030 - val_loss: 0.6115 - val_accuracy: 0.7300
Epoch 12/100
63/63 [=====] - 0s 4ms/step - loss: 0.2651 - accuracy: 0.8975 - val_loss: 0.5783 - val_accuracy: 0.7300
Epoch 13/100
63/63 [=====] - 0s 5ms/step - loss: 0.2415 - accuracy: 0.9035 - val_loss: 0.5673 - val_accuracy: 0.7320
Epoch 14/100
63/63 [=====] - 0s 4ms/step - loss: 0.2511 - accuracy: 0.8980 - val_loss: 0.5311 - val_accuracy: 0.7300
Epoch 15/100
63/63 [=====] - 0s 4ms/step - loss: 0.2662 - accuracy: 0.9105 - val_loss: 0.4920 - val_accuracy: 0.7520
Epoch 16/100
63/63 [=====] - 0s 5ms/step - loss: 0.2456 - accuracy: 0.9095 - val_loss: 0.5995 - val_accuracy: 0.7250
```



# Scoring & Plotting

score[0] will be test loss and score[1] will be test accuracy

```
score1 = cnn.evaluate(test_data, test_labels, verbose=0)
score2 = dense.evaluate(test_data, test_labels, verbose=0)
print(score1)
print(score2)
```

```
[1.0759947299957275, 0.8525000214576721]
[0.628600537776947, 0.7250000238418579]
```

**Green dotted line is where our model worked best**





+

4




+

# Experimenting With and Augmenting Field Data

# Field Data:

Data that is **different** from the **data that you used to build your model** (in your training and test sets). In the context of pneumonia imaging, an example of this could be images that were captured on a different x-ray than the images you built your model with.





# 73% Accuracy

```
y = (cnn.predict(field_data) > 0.5).astype("int32")  
print(accuracy_score(field_labels,y))
```

0.73

# Comparing and Contrasting Data


Function = `plot_one_image`

Field Data:


Varied in levels of blurriness

Training Data:

Consistent in levels of blurriness



# Solution: Augmentation



Accuracy  
Improved  
to 80%